

CS162 Discussion #7: Can You Find My House if I Give You My Virtual Address?

Announcements

- Phase 2 initial design due at 11:59pm – use `submit proj2-initial-design`
 - Design reviews Monday-Tuesday (10/18-10/19) in 258 Cory / 310 Soda: sign up!
 - Drop off a stapled hardcopy of your doc in the bag outside 346 Soda by tonight
 - Note: The comments you get during design review and on your doc are just comments...not a complete evaluation of your solution's correctness!
- Midterm 1 on Monday, 10/18 @ 6-9pm in 155 Dwinelle
 - Extra OH today @ 4-5pm in 651 Soda
 - Review session Sunday 10/17 @ 7-9pm in 306 Soda
- Phase 2 code due Tuesday 10/26 @ 11:59pm
 - Nachos Night Phase 2 will be Thursday 10/21 @ 7pm
 - `proj2-test` autograder is currently running every 4 hours

I'll page you ;)

Last week we talked about segmentation to create a virtual address space, but one of the disadvantages to simple segmentation (where each segment is a chunk of contiguous memory) is that variable segment sizes can cause memory to become fragmented. We can fix this problem by making sure that memory is always allocated in fixed-size chunks, which we call pages. These pages can then be assigned to different segments. Each page table entry contains a bunch of relevant information, including the physical page it's mapped to and permissions (read, write, valid, etc.). These pages are stored in a process' page table, which can be found by following the page table pointer. In a simple page table lookup, a process can locate the desired page based on virtual page number, then use the information stored in the page table entry to determine what physical page it should go to.

Not only that, we can also create segments out of pages. In this case, each virtual address will have a segment number in addition to a page number and an offset, and the physical address will be looked up using a multi-level page table. This page table will have one level indexed by segment number, which will contain another pointer to a page table indexed by virtual page number, and this entry will have the physical page number we're looking for. (We can also do the same thing just using two page indexes, without associating them with segment numbers.) This lets us conveniently group pages into sets, which is particularly useful in sparse memory spaces where there are groups of pages that aren't currently needed. Note that there may be permission bits at every level of the table, so we can check at every step of the translation process whether or not a page is valid, writeable, etc.

Trying to access pages that aren't currently in memory will cause a page fault, which occurs if a page table entry is marked as being invalid. The fault handler will then take care of bringing the page in from disk – we call this system demand paging. You'll see more of this during lecture in the coming week.

Cache Creek has fast access...to your cash D:

All that looking up in the page table seems like a lot of trouble, especially if you're going to be accessing the same virtual address a lot. It would be much better if we could remember where a virtual address took us the first time so we wouldn't have to go through that again – and that's what a cache is for. A cache keeps entries that are fast to look up, so we want to try to keep things we know will be looked up again in there to avoid taking the time to do the long address translation process all over again. The cache that's specially used for address translation is called the TLB.

Before it does a full-blown page table lookup, a process will check the TLB first to see if there's an entry already there that it can use. If not, this is a cache miss, and we don't have any choice but to go do the entire lookup. There are four different kinds of cache misses:

- *Compulsory*. The block hasn't been accessed before, so there's no way it could be in the cache.
- *Capacity*. There isn't enough room in the cache to keep all the blocks the process is trying to use.
- *Conflict*. Multiple memory locations are being mapped to the same place in the cache, so they kick each other out each time they're stored.
- *Coherence*. Memory has been updated so the entry isn't correct anymore.

In general, we do cache lookups by using an index to look for entries in the cache that might be what we're looking for. A tag then identifies the exact item we want. If nothing matches the tag, it's a cache miss. The specific lookup, however, depends on the type of cache we have. Here are some different kinds of caches:

- *Direct Mapped*. Each item is associated with one specific place in the cache using tag, index, and byte select.
- *N-way Set Associative*. Each index identifies a set of N blocks. The tag is compared between the N blocks to find the correct one, and then a byte select is used.
- *Fully Associative*. Every item can go in any block in the cache. The tag is compared between all blocks, then a byte select is used. There is no cache index here.

A TLB miss, which happens if the translation the process is looking for isn't currently stored in the TLB, will force us to look through the page table to find the page we want. Remember that pages don't necessarily have to be memory, though! Looking through the page table at this point could still cause a page fault, just like before. When the page has been found and loaded in memory, the TLB is updated with the new entry.

It's important to be careful how entries are being stored in the TLB, because accesses that cause a lot of conflicts can lead to thrashing. This means that you can potentially continuously create lots of cache misses if memory accesses are frequently kicking entries out of the TLB that the process will want to use again soon.